

Identifying Tree Species via Photographs of Leaves

Geoffrey E. Schneider

May 26, 2017

1 Introduction

My primary objective in this project is to practice, learn and experiment with machine learning techniques, as well as to form a kind of portfolio of my personal knowledge and abilities in this area. The project objective is to identify tree species from photos of their leaves. Because my objective is not the same as the project objective, I will not always take the most direct approach to the problem, experimenting with techniques that might not be absolutely necessary. I will however, try to mimic good practices in developing machine learning algorithms. As such, I will start with the easiest approach to the easiest version of the problem as outlined below in Section ??, as a baseline from which to compare other algorithms.

This is a work in progress, as apparent by the empty sections below.

2 First Approach: Logistic Regression on Two Species

I begin by using the easiest possible approach to get results for comparison with later attempts. The first approach attempts to classify pictures of leaves from two species: *abies concolor* and *abies nordmanniana* using logistic regression.

2.1 Data

There were a total of 86 sample pictures of these leaves taken from the Leafsnap database [1]. There were 51 for *abies concolor* and 35 for *abies nordmanniana* (See Figure 1 for examples). They were color jpg photos.

2.2 Preprocessing

The Octave script 'loadDataScript.m' loads and preprocesses the images. Images were all approximately 600×800 pixels (± 3 pixels per dimension) when rotated such that the first dimension was shorter than the second. Because this would result in a very large feature vector, the size was reduced to 30×40 pixels for all images via the Octave command 'imresize' from the 'image' package. The 86 resulting $30 \times 40 \times 3$ matrices were unrolled into a 86×3600 whose rows were randomly permuted. The data was normalized, and then divided into matrices of training data (X), cross-validation data (X_{cv}) and test data (X_{test}). These are labelled with corresponding vectors y , y_{cv} and y_{test} . The label 0 corresponds to *abies concolor* and 1 to *abies nordmanniana*.



Figure 1: Example photos of *abies concolor* (left) and *abies nordmanniana* (right).

2.3 Model

The model trained was a logistic regression model with regularization. The hypothesis for a particular example is given by:

$$h_{\theta}(x) = \frac{1}{1 + e^{x \cdot \theta}}$$

Where x is a feature vector and θ is a vector of the model parameters.

The objective function is given by cross-entropy

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m (y_i \log(h_{\theta}(x_i)) + (1 - y_i) \log(1 - h_{\theta}(x_i))) + \frac{\lambda}{2m} \sum_{k=1}^n \theta_k^2$$

where y_i is the label and x_i is the feature vector of the i th sample.

2.4 Training

The model parameters were trained by minimizing cross-entropy using the Octave 'fminunc' function, with theta initialized as the 0 vector.

2.5 Evaluation of results

The table below shows cost results for the model trained for various choices of the regularization parameter λ (after training, cost was tested with $\lambda = 0$).

λ :	0	0.1	0.3	1
training :	1.3115e-004	0.0016223	0.0028836	0.0067081
cross-validation:	1.1418e-004	0.0023556	0.0015548	0.0019042

The pre-training cost was 0.69315 on both the training set and the cross-validation set. If we predict *abies concolor* when our hypothesis function gives a result < 0.5 , this gives a 100%



Figure 2: Visual representation of the trained parameters of the logistic regression model.

accuracy on the cross-validation set for all choices of λ . Figure 2 shows a visual representation of the model parameters (without the bias) when trained with $\lambda = 0$.

The 100 % accuracy clearly shows that the data is linearly separable, even with the large loss of information due to preprocessing. Because this simple model easily solves this problem, we must make the problem more difficult by including more species so that this can be a valid basis for comparison.

The algorithm was also tested on lab images from the leafsnap dataset in order to ensure that the result was not obtained due to systematic differences (e.g. lighting differences) between photos of the two species. The same 100 % accuracy held in this case.

2.6 Random Initialization

Using a random mean 0, variance 1 Gaussian initialization with $\lambda = 0$, we get an even lower cross-entropy of 1.4391e-006 on the cross-validation set. Figure 3 shows that the model parameters train to radically different values in this case.

2.7 More Species

To get a better basis for comparison, the logistic regression model was additionally trained to distinguish 10 and 100 species. A few examples of these species are shown in Figure 4. This was done by training 10 and 100 logistic regression models respectively, each distinguishing one species from any of the others. Because there were no overfitting issues in the 2 species problem, the regularization parameter was set to 0, and no other such parameters were considered (this was also done for time considerations). The results are shown below, reported as percent accuracy on




Figure 3: Visual representation of the trained parameters of the logistic regression model with random initialization.

the cross-validation set. Random initialization was Gaussian with mean 0 and variance 0.01.

# species :	2	10	100
Zero initialization :	100%	81.8%	49.6%
Random initialization:	100%	84.1%	48.8%

As before, the algorithm was also tested on the lab dataset (for 10 species) and an even stronger result (94.6 % accuracy) was obtained.

3 Implementing NetMorph

NetMorph [3] is an algorithm that takes a neural net and factors layers to obtain a deeper net (with activations that require training) that initially gives the same (or very similar) outputs to the original net. This allows us to use what previously trained models "know" to train new models. This part of the project was implemented in TensorFlow with the Python API.

References

- [1] Neeraj Kumar, Peter N. Belhumeur, Arijit Biswas, David W. Jacobs, W. John Kress, Ida Lopez, João V. B. Soares, "Leafsnap: A Computer Vision System for Automatic Plant Species Identification," Proceedings of the 12th European Conference on Computer Vision (ECCV), October 2012.

[2]



Figure 4: Example photos of *acer palmatum* (upper left), *acer pseudoplatanus* (upper right), *betula nigra* (bottom left) and *crataegus viridis* (bottom right). The two species on top are included in the models trained on 10 species and all four are included in the models trained on 100 species.




Figure 5: Learning for 200 iterations with the Adam Optimizer [2]. Factoring a 2×2 matrix of 3×3 filters into two 2×2 matrices of 2×2 filters.




Figure 6: Learning iterations 200 to 5000 with the Adam Optimizer.

